

Caret

Practical Machine Learning (with R)

CARET



Caret

“Misc functions for training and plotting classification and regression models.”

⇒ Really:

- Wraps 100's of modeling functions
- Automates tediousness of model building
- Manages a process

⇒ Competitors:

- [Rattle](#) : Graham Williams et al./Togaware.com
- [R Commander](#) : Statistical workbench



Caret Goals

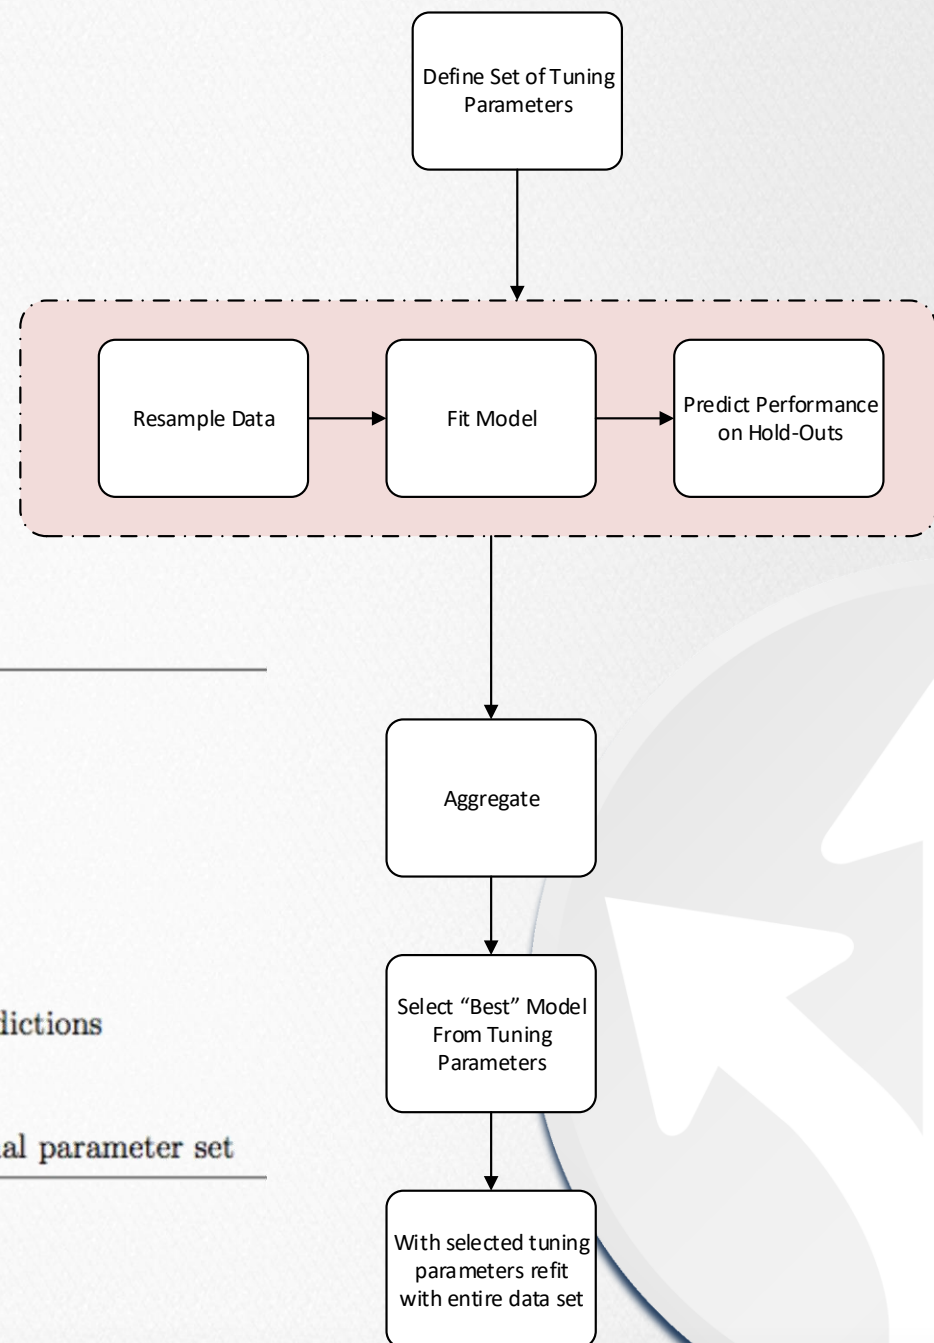
Does a couple things:

- Preprocess data (transforms, imputes)
- evaluate, using resampling, the effect of model tuning parameters on performance
- choose the “optimal” model across these parameters
- estimate model performance from a training set
- Variable Importance
- Aids feature selection



Process

```
1 Define sets of model parameter values to evaluate
2 for each parameter set do
3   for each resampling iteration do
4     Hold-out specific samples
5     [Optional] Pre-process the data
6     Fit the model on the remainder
7     Predict the hold-out samples
8   end
9   Calculate the average performance across hold-out predictions
10 end
11 Determine the optimal parameter set
12 Fit the final model to all the training data using the optimal parameter set
```



LOTS OF CONFIGURATIONS

- ⇒ Easy if you know what you are doing
- ⇒ which method?

Caret Model List*

- ⇒ Controlled mostly through
 - `train (tuneLength, tuneGrid)`
 - `trainControl` – supplied to `train`

